
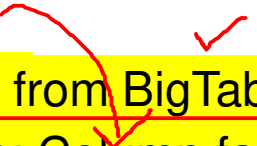








- Google BigTable google file system
 - High performance data storage system built on GFS and other Google technologies
 - Master-slave architecture
 - One key, multiple values
 - ✓ Columnar, SSTable (Sorted String Table) Storage, Append-only, Memtable, Compaction
- Amazon DynamoDb 
 - Highly available and scalable key-value storage system
 - Decentralized peer to peer architecture
 - Compromise on consistency for better availability - Eventual consistency
 - Consistent hashing, Gossip protocol, ~~Replication~~, Read repair
- Cassandra 
 - Inherited from BigTable and DynamoDb
 - BigTable: Column families, Memtable, SSTable
 - DynamoDb: Consistent hashing, Partitioning, Replication 





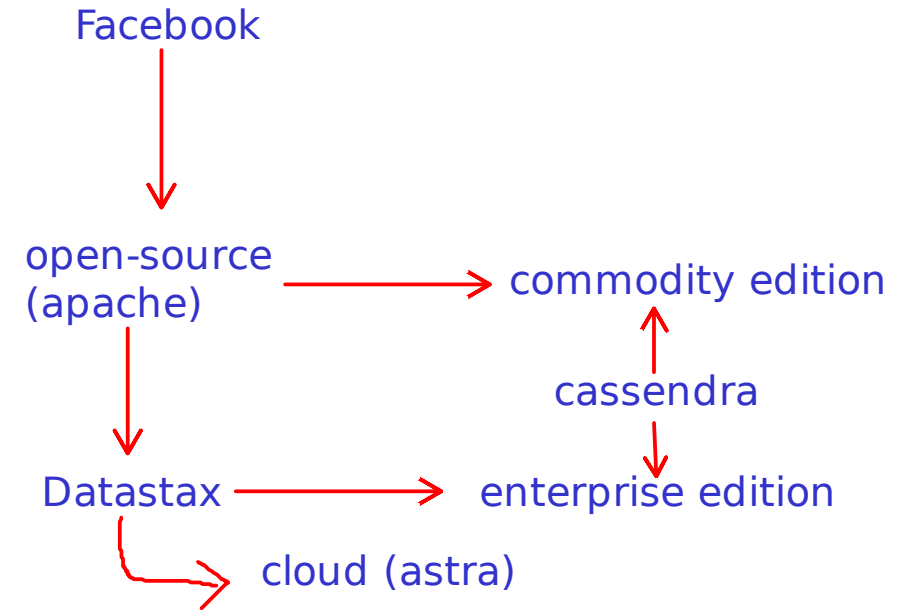
	ename	job	sal
1	abc	1 ck	1 1202
2	pqr	2 mgr	2 2200

row id
row key



Introduction

- **Developed by**
 - Avinash Laxman (Co-inventor Amazon DynamoDb)
 - Prashant Malik (Technical Leader at Facebook) ✓
- **Goals:**
 - Distributed NoSQL database (on commodity hardware)
 - Large amount of structured data
 - High availability
 - No single point of failure
- Basic data model is rows & columns
- Column-oriented, Decentralized peer to peer & follow Eventual consistency
- Datastax company develop and support commercial edition of Cassandra



Cassandra Development

- Developed in Java ✓
- 2007-2008 - Developed at Facebook ✓
- July 2008 - Open sourced by Facebook ✓
- March 2009 - Apache Incubator project ✓
- February 2010 - Apache Top-level project ✓
- 2011 - version 0.8 - Added CQL ✓
- 2013 - version 2.0 - Added light-weight transactions, Triggers ✓
- 2015 - version 3.0 - Storage engine improved, Materialized views ✓
- 2017 -version 3.11 – bug fix from the last release ✓
- 2021 - version 4.0.5 ---> we r using this version ✓
- 2022 - version 4.1.0 - Latest release



Installation

- Prerequisite
 - Java 8 (Java 11 experimental)
- Can be installed through apt or yum tool (Ubuntu/CentOS)
- Manual installation
 - Download Cassandra 3.11.x (.tar.gz) and extract it
 - set CASSANDRA_HOME to Cassandra directory
 - set JAVA_HOME to JDK 8 directory
 - Install python 2.7 (for cqlsh)
 - set CASSANDRA_HOME/bin into PATH variable
 - Start Cassandra
- terminal1> cassandra
- terminal2> cqlsh

→ mapping

mysql	cassandra	
-----	-----	
database	keyspace	
table	table	
row	row	
column	column	



Features

- Peer to peer architecture no master-slave SPOF
- Linear scale performance capacity \propto nodes
- High Performance \longrightarrow high speed in read/write
- Simplified deployment and maintenance \rightarrow linux
- Less expensive \Rightarrow horiz scalability
- Supports multiple programming languages \longrightarrow py , java ,c++ , Node.js, C#....., &REST service
- Operational and Development simplicity
- Cloud Availability AWS, Azur,GCP
- Ability to deploy across data centers
- Fault tolerant
- Configurable consistency (tight or eventual)
- Flexible data model
- Column family store



Limitations

- Aggregation operations are not supported
- Range queries on partition key are not supported
- Not good for too many joins
- Not suitable for transactional data ✓
- During compaction performance / throughput slows down
- Not designed for update-delete



Performance

- Performance measures
 - Throughput (operations per second) ↑
 - Latency (time required for one operation) ↓
- Cassandra vs MySQL
 - MySQL (more than 50GB data)
 - Write speed: 300 ms ✓
 - Read speed: 350 ms ✓
 - Cassandra (more than 50GB data)
 - Write speed: 0.12 ms ←
 - Read speed: 15 ms ←




- Applications

- Product catalog/Playlist ✓
- Recommendation/Personalization engine ✓
- Sensor/IoT data ✓
- Messaging/Time-series data ✓
- Fraud detection

- Customers

- Facebook, Netflix, eBay, Apple, Walmart, GoDaddy

- Application requirements

- Store and handle time-series data ✓
 - Store and handle large volume of data ✓
 - Scale predictably (Linear Scaling) ✓
 - High availability ✓
- 

uuid -> universally unique identifier
128bit -> 16byte

time-uuid = time stamp + uuid
(helpful in sorting
on time field)



Data Model

- Cassandra provides the Cassandra Query Language (**CQL**), an **SQL-like** language, to create and update database schema and access data
- CQL allows users to organize data within a cluster of Cassandra nodes using:
 - **Keyspace**
 - Defines how a dataset is replicated, per datacenter
 - Replication is the number of copies saved per cluster. **Keyspaces contain tables**
 - **Table**
 - Defines the typed schema for a collection of partitions. Tables contain partitions, which contain rows, which contain columns. **Cassandra tables can flexibly add new columns to tables with zero downtime.**
 - **Partition**
 - Defines the mandatory part of the primary key all rows in Cassandra must have to identify the node in a cluster where the row is stored
 - **All performant queries supply the partition key in the query.**
 - **Row**
 - Contains a collection of columns identified by a unique primary key made up of the partition key and optionally **additional clustering keys**
 - **Column**
 - A single datum with a type which belongs to a row



- Users can access Cassandra through its nodes using Cassandra Query Language (CQL)
- CQL treats the database (Keyspace) as a container of tables
- Programmers use cqlsh: a prompt to work with CQL or separate application language drivers



Data Types

- ✓ ascii: US-ascii character string
- ✓ bigint: 64-bit signed long ints
- ✓ blob: Arbitrary bytes in hexadecimal
- ✓ boolean: True or False
- ✓ counter: Distributed counter values 64 bit
- ✓ decimal: Variable precision decimal
- ✓ double: 64-bit floating point
- ✓ float: 32-bit floating point
- ✓ frozen: Tuples, collections, UDT containing CQL types
- ✓ inet - IP address in ipv4 or ipv6 string format
- ✓ int: 32 bit signed integer
- ✓ list: Collection of elements
- ✓ map: JSON style collection of elements
- ✓ set: Sorted collection of elements
- ✓ text: UTF-8 encoded strings
- ✓ timestamp: ID generated with
- ✓ date+time: as int/string
- ✓ timeuuid: Type 1 uuid
- ✓ tuple: A group of 2,3 fields
- ✓ uuid: Standard uuid (128-bit)
- ✓ varchar: UTF-8 encoded string
- ✓ varint: Arbitrary precision integer



Cassandra vs MongoDB: Differences

■ Cassandra

- Dev In java ✓
- Column based ✓
- Cassandra uses a traditional model with a table structure, using rows and columns. ✓
- Cassandra offers an assortment of peers node
- Used in-house query language, CQL ✓
- no internal aggregation framework ✓
- According to the CAP theorem Cassandra is an AP system.

■ MongoDB

- Dev in C++ ✓
- Document based ✓
- MongoDB employs an objective-oriented or data-oriented model
- MongoDB uses a single master node.
- queries are structured into JSON fragments
- own aggregation framework ✓
- According to the CAP theorem, MongoDB is a CP system

